

Execute the code with the following command

Result of the execution

(The response you will get may be slightly different in wording from what you see below)

We will cover Memory in more detail in future tutorials. You can find API documentation for Memory [here](#)

We hope this Step-By-Step Tutorial was helpful in getting you started with LangChain.js. We will be digging deeper into the individual modules and use cases in upcoming tutorials. We can't wait. We love LangChain. **STAY TUNED!**



Sign Up

Sign In

🔍 Search packages

Search

openai **TS**

3.2.1 • Public • Published a month ago

 Readme

 Code Beta

 2 Dependencies

 437 Dependents

 25 Versions

OpenAI Node.js Library

The OpenAI Node.js library provides convenient access to the OpenAI API from Node.js applications. Most of the code in this library is generated from our **OpenAPI specification**.

Important note: this library is meant for server-side usage only, as using it in client-side browser code will expose your secret API key. [See here](#) for more details.

Installation

```
$ npm install openai
```

Usage

The library needs to be configured with your account's secret key, which is available on the [website](#). We recommend setting it as an environment variable. Here's an example of

initializing the library with the API key loaded from an environment variable and creating a completion:

```
const { Configuration, OpenAIApi } = require("openai");

const configuration = new Configuration({
  apiKey: process.env.OPENAI_API_KEY,
});
const openai = new OpenAIApi(configuration);

const completion = await openai.createCompletion({
  model: "text-davinci-003",
  prompt: "Hello world",
});
console.log(completion.data.choices[0].text);
```

Check out the [full API documentation](#) for examples of all the available functions.

Request options

All of the available API request functions additionally contain an optional final parameter where you can pass custom **axios request options**, for example:

```
const completion = await openai.createCompletion(
  {
    model: "text-davinci-003",
    prompt: "Hello world",
  },
  {
    timeout: 1000,
    headers: {
      "Example-Header": "example",
    },
  }
);
```

Error handling

API requests can potentially return errors due to invalid inputs or other issues. These errors can be handled with a `try...catch` statement, and the error details can be found in either `error.response` or `error.message`:

```
try {
  const completion = await openai.createCompletion({
    model: "text-davinci-003",
    prompt: "Hello world",
  });
  console.log(completion.data.choices[0].text);
} catch (error) {
  if (error.response) {
    console.log(error.response.status);
    console.log(error.response.data);
  } else {
    console.log(error.message);
  }
}
```

Streaming completions

Streaming completions (`stream=true`) are not natively supported in this package yet, but **a workaround exists** if needed.

Upgrade guide

All breaking changes for major version releases are listed below.

3.0.0

- The function signature of `createCompletion(engineId, params)` changed to `createCompletion(params)`. The value previously passed in as the `engineId` argument should now be passed in as `model` in the `params` object (e.g. `createCompletion({ model: "text-davinci-003", ... })`)
- Replace any `createCompletionFromModel(params)` calls with `createCompletion(params)`

Thanks

Thank you to **ceifa** for creating and maintaining the original unofficial `openai` npm package before we released this official library! ceifa's original package has been renamed to **gpt-x**.

Keywords

openai open ai gpt-3 gpt3

Install

```
> npm i openai
```

Repository

 github.com/openai/openai-node

Homepage

 github.com/openai/openai-node#readme

↓ Weekly Downloads

449,234

Version

3.2.1

License

MIT

Unpacked Size

480 kB

Total Files

26

Last publish

a month ago

Collaborators



[>Try on RunKit](#)

[🚩Report malware](#)



Support

[Help](#)

[Advisories](#)

[Status](#)

[Contact npm](#)

Company

[About](#)

[Blog](#)

[Press](#)

Terms & Policies

[Policies](#)

[Terms of Use](#)

[Code of Conduct](#)

[Privacy](#)





Node.JS Client



This page provides installation instructions, usage examples, and a reference for the [Pinecone Node.JS client](#).

Warning

This is a **public preview** ("Beta") client. Test thoroughly before using this client for production workloads. No SLAs or technical support commitments are provided for this client. Expect potential breaking changes in future releases.

Getting Started

Installation

Use the following shell command to install the Node.JS client for use with Node.JS versions 17 and above:

Shell

```
npm install @pinecone-database/pinecone
```

Alternatively, you can install Pinecone with Yarn:

Shell

```
yarn add @pinecone-database/pinecone
```

Usage

Initialize the client

To initialize the client, instantiate the `PineconeClient` class and call the `init` method. The `init` method takes an object with the `apiKey` and `environment` properties:

JavaScript

```
import { PineconeClient } from "@pinecone-database/pinecone";

const pinecone = new PineconeClient();
await pinecone.init({
  environment: "YOUR_ENVIRONMENT",
  apiKey: "YOUR_API_KEY",
});
```

Create index

The following example creates an index without a metadata configuration. By default, Pinecone indexes all metadata.

JavaScript

```
await pinecone.createIndex({
  createRequest: {
    name: "example-index",
    dimension: 1024,
  },
});
```

The following example creates an index that only indexes the "color" metadata field. Queries against this index cannot filter based on any other metadata field.

JavaScript

```
await pinecone.createIndex({
  createRequest: {
    name: "example-index-2",
    dimension: 1024,
    metadata_config: {
      indexed: ["color"],
    },
  },
});
```

List indexes

The following example logs all indexes in your project.

JavaScript

```
const indexesList = await pinecone.listIndexes();
```

Describe index

The following example logs information about the index `example-index`.

JavaScript

```
const indexDescription = await pinecone.describeIndex({
  indexName: "example-index",
});
```

Delete index

The following example deletes `example-index`.

JavaScript

```
await pinecone.deleteIndex({
  indexName: "example-index",
});
```

Scale replicas

The following example sets the number of replicas and pod type for `example-index`.

JavaScript

```
await pinecone.configureIndex({
  indexName: "example-index",
  patchRequest: {
    replicas: 2,
    podType: "p2",
  },
});
```

Describe index statistics

The following example returns statistics about the index `example-index`.

JavaScript

```
const index = pinecone.Index("example-index");
const indexStats = index.describeIndexStats({
  describeIndexStatsRequest: {
    filter: {},
  },
});
```

```
});
```

Upsert vectors

The following example upserts vectors to `example-index`.

JavaScript

```
const index = pinecone.Index("example-index");
const upsertRequest = {
  vectors: [
    {
      id: "vec1",
      values: [0.1, 0.2, 0.3, 0.4],
      metadata: {
        genre: "drama",
      },
    },
    {
      id: "vec2",
      values: [0.2, 0.3, 0.4, 0.5],
      metadata: {
        genre: "action",
      },
    },
  ],
  namespace: "example-namespace",
};
const upsertResponse = await index.upsert({ upsertRequest });
```

Query an index

The following example queries the index `example-index` with metadata filtering.

JavaScript

```
const index = pinecone.Index("example-index");
const queryRequest = {
  vector: [0.1, 0.2, 0.3, 0.4],
  topK: 10,
  includeValues: true,
  includeMetadata: true,
  filter: {
    genre: { $in: ["comedy", "documentary", "drama"] },
  },
  namespace: "example-namespace",
};
const queryResponse = await index.query({ queryRequest });
```

Delete vectors

The following example deletes vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");
await index.delete({
  ids: ["vec1", "vec2"],
  namespace: "example-namespace",
});
```

Fetch vectors

The following example fetches vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");
const fetchResponse = await index.fetch({
  ids: ["vec1", "vec2"],
  namespace: "example-namespace",
});
```

Update vectors

The following example updates vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");
const updateRequest = {
  id: "vec1",
  values: [0.1, 0.2, 0.3, 0.4],
  setMetadata: { genre: "drama" },
  namespace: "example-namespace",
};
const updateResponse = await index.update({ updateRequest });
```

Create collection

The following example creates the collection `example-collection` from `example-index`.

JavaScript

```
const createCollectionRequest = {
  name: "example-collection",
  source: "example-index",
};
```

```
};

await pinecone.createCollection({
  createCollectionRequest,
});
```

List collections

The following example returns a list of the collections in the current project.

JavaScript

```
const collectionsList = await pinecone.listCollections();
```

Describe a collection

The following example returns a description of the collection `example-collection`.

JavaScript

```
const collectionDescription = await pinecone.describeCollection({
  collectionName: "example-collection",
});
```

Delete a collection

The following example deletes the collection `example-collection`.

JavaScript

```
await pinecone.deleteCollection({
  collectionName: "example-collection",
});
```

Reference

For the REST API or other clients, see [the API reference](#).

init()

```
pinecone.init(configuration: PineconeClientConfiguration)
```

Initialize the Pinecone client.

Parameters	Type	Description
<code>configuration</code>	<code>PineconeClientConfiguration</code>	The configuration for the Pinecone client.

Types

`PineconeClientConfiguration`

Parameters	Type	Description
<code>apiKey</code>	<code>string</code>	The API key for the Pinecone service.
<code>environment</code>	<code>string</code>	The cloud environment of your Pinecone project.

Example:

JavaScript

```
import { PineconeClient } from "@pinecone-database/pinecone";
const pinecone = new PineconeClient();
await pinecone.init({
  apiKey: "YOUR_API_KEY",
  environment: "YOUR_ENVIRONMENT",
});
```

`configureIndex()`

```
pinecone.configure_index(indexName: string, patchRequest?: PatchRequest)
```

Configure an index to change pod type and number of replicas.

Parameters	Type	Description
<code>requestParameters</code>	<code>ConfigureIndexRequest</code>	Index configuration parameters.

Types

`ConfigureIndexRequest`

Parameters	Type	Description
<code>indexName</code>	<code>string</code>	The name of the index.
<code>patchRequest</code>	<code>PatchRequest</code>	(Optional) Patch request parameters.

`PatchRequest`

Parameters	Type	Description
<code>replicas</code>	number	(Optional) The number of replicas to configure for this index.
<code>podType</code>	string	(Optional) The new pod type for the index. One of <code>s1</code> , <code>p1</code> , or <code>p2</code> appended with <code>.</code> and one of <code>x1</code> , <code>x2</code> , <code>x4</code> , or <code>x8</code> .

Example:

JavaScript

```
const newNumberOfReplicas = 4;
const newPodType = "s1.x4";
await pinecone.configureIndex({
  indexName: "example-index",
  patchRequest: {
    replicas: newNumberOfReplicas,
    podType: newPodType,
  },
});
```

createCollection()

```
pinecone.createCollection(requestParameters: CreateCollectionOperationRequest)
```

Create a collection from an index.

Parameters	Type	Description
<code>requestParameters</code>	CreateCollectionOperationRequest	Create collection operation wrapper

Types

CreateCollectionOperationRequest

Parameters	Type	Description
<code>createCollectionRequest</code>	CreateCollectionRequest	Collection request parameters.

CreateCollectionRequest

Parameters	Type	Description
<code>name</code>	string	The name of the collection to be created.

Parameters	Type	Description
<code>source</code>	string	The name of the source index to be used as the source for the collection.

Example:

JavaScript

```
await pinecone.createCollection({
  createCollectionRequest: {
    name: "example-collection",
    source: "example-index",
  },
});
```

createIndex()

```
pinecone.createIndex(requestParameters?: CreateIndexRequest)
```

Create an index.

Parameters	Type	Description
<code>requestParameters</code>	CreateIndexRequest	Create index operation wrapper

Types

CreateIndexRequest

Parameters	Type	Description
<code>createRequest</code>	CreateRequest	Create index request parameters

CreateRequest

Parameters	Type	Description
<code>name</code>	str	The name of the index to be created. The maximum length is 45 characters.
<code>dimension</code>	integer	The dimensions of the vectors to be inserted in the index.
<code>metric</code>	str	(Optional) The distance metric to be used for similarity search: 'euclidean', 'cosine', or 'dotproduct'.

Parameters	Type	Description
pods	int	(Optional) The number of pods for the index to use, including replicas.
replicas	int	(Optional) The number of replicas.
pod_type	str	(Optional) The new pod type for the index. One of s1 , p1 , or p2 appended with . and one of x1 , x2 , x4 , or x8 .
metadata_config	object	(Optional) Configuration for the behavior of Pinecone's internal metadata index. By default, all metadata is indexed; when metadata_config is present, only specified metadata fields are indexed. To specify metadata fields to index, provide a JSON object of the following form: { "indexed": ["example_metadata_field"] }
source_collection	str	(Optional) The name of the collection to create an index from.

Example:

JavaScript

```
// The following example creates an index without a metadata
// configuration. By default, Pinecone indexes all metadata.
await pinecone.createIndex({
  createRequest: {
    name: "pinecone-index",
    dimension: 1024,
  },
});

// The following example creates an index that only indexes
// the 'color' metadata field. Queries against this index
// cannot filter based on any other metadata field.

await pinecone.createIndex({
  createRequest: {
    name: "example-index-2",
    dimension: 1024,
    metadata_config: {
      indexed: ["color"],
    },
  },
});
```

deleteCollection()

```
pinecone.deleteCollection(requestParameters: DeleteCollectionRequest)
```

Delete an existing collection.

Types

Parameters	Type	Description
<code>requestParameters</code>	DeleteCollectionRequest	Delete collection request parameters

DeleteCollectionRequest

Parameters	Type	Description
<code>collectionName</code>	string	The name of the collection to delete.

Example:

JavaScript

```
await pinecone.deleteCollection({
  collectionName: "example-collection",
});
```

deleteIndex()

```
pinecone.deleteIndex(requestParameters: DeleteIndexRequest)
```

Delete an index.

Types

Parameters	Type	Description
<code>requestParameters</code>	DeleteIndexRequest	Delete index request parameters

DeleteIndexRequest

Parameters	Type	Description
<code>indexName</code>	string	The name of the index to delete.

Example:

JavaScript

```
await pinecone.deleteIndex({
  indexName: "example-index",
});
```

describeCollection()

```
pinecone.describeCollection(requestParameters: DescribeCollectionRequest)
```

Get a description of a collection.

Types

Parameters	Type	Description
<code>requestParameters</code>	<code>DescribeCollectionRequest</code>	Describe collection request parameters

DescribeCollectionRequest

Parameters	Type	Description
<code>collectionName</code>	<code>string</code>	The name of the collection.

Example:

JavaScript

```
const collectionDescription = await pinecone.describeCollection({
  collectionName: "example-collection",
});
```

Return:

- `collectionMeta` : `object` Configuration information and deployment status of the collection.
 - `name` : `string` The name of the collection.
 - `size` : `integer` The size of the collection in bytes.
 - `status` : `string` The status of the collection.

describeIndex()

```
pinecone.describeIndex(requestParameters: DescribeIndexRequest)
```

Get a description of an index.

Types

Parameters	Type	Description
<code>requestParameters</code>	<code>DescribeIndexRequest</code>	Describe index request parameters

DescribeIndexRequest

Parameters	Type	Description
<code>indexName</code>	<code>string</code>	The name of the index.

Types

Returns:

- `database` : `object`
- `name` : `string` The name of the index.
- `dimension` : `integer` The dimensions of the vectors to be inserted in the index.
- `metric` : `string` The distance metric used for similarity search: 'euclidean', 'cosine', or 'dotproduct'.
- `pod` : `integer` The number of pods the index uses, including replicas.
- `replicas` : `integer` The number of replicas.
- `pod_type` : `string` The pod type for the index. One of `s1` , `p1` , or `p2` appended with `.` and one of `x1` , `x2` , `x4` , or `x8` .
- `metadata_config` : `object` Configuration for the behavior of Pinecone's internal metadata index. By default, all metadata is indexed; when `metadata_config` is present, only specified metadata fields are indexed. To specify metadata fields to index, provide a JSON object of the following form: `{"indexed": ["example_metadata_field"]}`
- `status` : `object`
- `ready` : `boolean` Whether the index is ready to serve queries.
- `state` : `string` One of `Initializing` , `ScalingUp` , `ScalingDown` , `Terminating` , or `Ready` .

Example:

JavaScript

```
const indexDescription = await pinecone.describeIndex({
  indexName: "example-index",
});
```

listCollections

```
pinecone.listCollections()
```

Return a list of the collections in your project.

Example:

```
JavaScript
```

```
const collections = await pinecone.listCollections();
```

Returns:

- `array` of `strings` The names of the collections in your project.

listIndexes

```
pinecone.listIndexes()
```

Return a list of your Pinecone indexes.

Returns:

- `array` of `strings` The names of the indexes in your project.

Example:

```
JavaScript
```

```
const indexesList = await pinecone.listIndexes();
```

Index()

```
pinecone.Index(indexName: string)
```

Construct an Index object.

Parameters	Type	Description
<code>indexName</code>	<code>string</code>	The name of the index.

Example:

```
JavaScript
```

```
const index = pinecone.Index("example-index");
```

Index.delete1()

```
index.delete(requestParameters: Delete1Request)
```

Delete items by their ID from a single namespace.

Parameters	Type	Description
<code>requestParameters</code>	Delete1Request	Delete request parameters

Types

Delete1Request

Parameters	Type	Description
<code>ids</code>	Array	(Optional) The IDs of the items to delete.
<code>deleteAll</code>	boolean	(Optional) Indicates that all vectors in the index namespace should be deleted.
<code>namespace</code>	str	(Optional) The namespace to delete vectors from, if applicable.

Types

Example:

JavaScript

```
await index.delete1({
  ids: ["example-id-1", "example-id-2"],
  namespace: "example-namespace",
});
```

Index.describeIndexStats()

```
index.describeIndexStats(requestParameters: DescribeIndexStatsOperationRequest)
```

Returns statistics about the index's contents, including the vector count per namespace and the number of dimensions.

Parameters	Type	Description
<code>requestParameters</code>	DescribeIndexStatsOperationRequest	Describe index stats request wrapper

Types

DescribeIndexStatsOperationRequest

Parameters	Type	Description
<code>describeIndexStatsRequest</code>	<code>DescribeIndexStatsRequest</code>	Describe index stats request parameters

DescribeIndexStatsRequest

parameter	Type	Description
<code>filter</code>	<code>object</code>	(Optional) A metadata filter expression.

Returns:

- `namespaces` : `object` A mapping for each namespace in the index from the namespace name to a summary of its contents. If a metadata filter expression is present, the summary will reflect only vectors matching that expression.
- `dimension` : `int64` The dimension of the indexed vectors.
- `indexFullness` : `float` The fullness of the index, regardless of whether a metadata filter expression was passed. The granularity of this metric is 10%.
- `totalVectorCount` : `int64` The total number of vectors in the index.

Example:

JavaScript

```
const indexStats = await index.describeIndexStats({
  describeIndexStatsRequest: {},
});
```

Read more about [filtering](#) for more detail.

Index.fetch()

```
index.fetch(requestParameters: FetchRequest)
```

The Fetch operation looks up and returns vectors, by ID, from a single namespace. The returned vectors include the vector data and metadata.

Parameters	Type	Description
<code>requestParameters</code>	<code>FetchRequest</code>	Fetch request parameters

Types

FetchRequest

Parameters	Type	Description
<code>ids</code>	Array	The vector IDs to fetch. Does not accept values containing spaces.
<code>namespace</code>	string	(Optional) The namespace containing the vectors.

Returns:

- `vectors` : `object` Contains the vectors.
- `namespace` : `string` The namespace of the vectors.

Example:

JavaScript

```
const fetchResponse = await index.fetch({  
  ids: ["example-id-1", "example-id-2"],  
  namespace: "example-namespace",  
});
```

Index.query()

```
index.query(requestParameters: QueryOperationRequest)
```

Search a namespace using a query vector. Retrieves the ids of the most similar items in a namespace, along with their similarity scores.

Parameters	Type	Description
<code>requestParameters</code>	QueryOperationRequest	The query operation request wrapper.

Types

Parameters	Type	Description
<code>queryRequest</code>	QueryRequest	The query operation request.

QueryRequest

Parameter	Type	Description
<code>namespace</code>	string	(Optional) The namespace to query.

Parameter	Type	Description
<code>topK</code>	number	The number of results to return for each query.
<code>filter</code>	object	(Optional) The filter to apply. You can use vector metadata to limit your search. See https://www.pinecone.io/docs/metadata-filtering/ .
<code>includeValues</code>	boolean	(Optional) Indicates whether vector values are included in the response. Defaults to <code>false</code> .
<code>includeMetadata</code>	boolean	(Optional) Indicates whether metadata is included in the response as well as the ids. Defaults to <code>false</code> .
<code>vector</code>	Array	(Optional) The query vector. This should be the same length as the dimension of the index being queried. Each <code>query()</code> request can contain only one of the parameters <code>id</code> or <code>vector</code> .
<code>id</code>	string	(Optional) The unique ID of the vector to be used as a query vector. Each <code>query()</code> request can contain only one of the parameters <code>vector</code> or <code>id</code> .

Example:

JavaScript

```
const queryResponse = await index.query({
  queryRequest: {
    namespace: "example-namespace",
    topK: 10,
    filter: {
      genre: { $in: ["comedy", "documentary", "drama"] },
    },
    includeValues: true,
    includeMetadata: true,
    vector: [0.1, 0.2, 0.3, 0.4],
  },
});
```

Index.update()

```
index.update(requestParameters: UpdateOperationRequest)
```

Updates vectors in a namespace. If a value is included, it will overwrite the previous value. If `setMetadata` is included in the `updateRequest`, the values of the fields specified in it will be added or overwrite the previous value.

Parameters	Type	Description
<code>requestParameters</code>	<code>UpdateOperationRequest</code>	The update operation wrapper

Types

UpdateOperationRequest

Parameters	Type	Description
<code>updateRequest</code>	<code>UpdateRequest</code>	The update request.

UpdateRequest

Parameter	Type	Description
<code>id</code>	string	The vector's unique ID.
<code>values</code>	Array	(Optional) Vector data.
<code>setMetadata</code>	object	(Optional) Metadata to set for the vector.
<code>namespace</code>	string	(Optional) The namespace containing the vector.

Example:

JavaScript

```
const updateResponse = await index.update({
  updatedRequest: {
    id: "vec1",
    values: [0.1, 0.2, 0.3, 0.4],
    setMetadata: {
      genre: "drama",
    },
    namespace: "example-namespace",
  },
});
```

Index.upsert()

```
index.upsert(requestParameters: UpsertOperationRequest)
```

Writes vectors into a namespace. If a new value is upserted for an existing vector ID, it will overwrite the previous value.

Parameters	Type	Description
<code>requestParameters</code>	<code>UpsertOperationRequest</code>	Upsert operation wrapper

Types

`UpsertOperationRequest`

Parameters	Type	Description
<code>upsertRequest</code>	<code>UpsertRequest</code>	The upsert request.

`UpsertRequest`

| Parameter | Type | Description |

| `vectors` | Array | An array containing the vectors to upsert. Recommended batch limit is 100 vectors.

| `id` (str) - The vector's unique id.

| `values` ([float]) - The vector data.

| `metadata` (object) - (Optional) Metadata for the vector. |

| `namespace` | string | (Optional) The namespace name to upsert vectors. |

Vector

Parameter	Type	Description
<code>id</code>	string	The vector's unique ID.
<code>values</code>	Array	Vector data.
<code>metadata</code>	object	(Optional) Metadata for the vector.

Returns:

- `upsertedCount` : `int64` The number of vectors upserted.

Example:

JavaScript

```
const upsertResponse = await index.upsert({
  upsertRequest: {
    vectors: [
      {
        id: "vec1",
        values: [0.1, 0.2, 0.3, 0.4],
        metadata: {
          genre: "drama",
        },
      },
    ],
  },
});
```

```
{
  id: "vec2",
  values: [0.1, 0.2, 0.3, 0.4],
  metadata: {
    genre: "comedy",
  },
},
],
namespace: "example-namespace",
},
});
```

 Updated 5 days ago



[Pricing](#) [Docs](#) [Learn](#) [Company](#) [Contact](#) [Careers](#)

© Pinecone Systems, Inc. | San Francisco, CA | [Terms](#) | [Privacy](#) | [Product Privacy](#) | [Cookies](#) | [Trust & Security](#) | [System Status](#)

Pinecone is a registered trademark of Pinecone Systems, Inc.



Sophia Yang

Apr 8 · 6 min read · Member-only · Listen



4 Ways to Do Question Answering in LangChain

Chat with your long PDF docs: load_qa_chain, RetrievalQA, VectorstoreIndexCreator, ConversationalRetrievalChain

Are you interested in chatting with your own documents, whether it is a text file, a PDF, or a website? LangChain makes it easy for you to do question answering with your documents. But do you know that there are at least 4 ways to do question answering in LangChain? In this blog post, we are going to explore four different ways to do question-answering and the various options you could consider for your use cases.

Before we dive into question answering, you may wonder: what is LangChain? Great question! In my opinion, LangChain is the easiest way to interact with language models and build applications. It is an **open-source** tool that wraps around many LLMs and tools. Check out my previous [blog post](#) and [video](#) on an overview of LangChain.

Okay, now let's get started with question-answering on external documents.

Code: Check out the code for this blog post [here](#).

Set up OpenAI API

Create an account at OpenAI and create an API key: <https://platform.openai.com/account>. Note that OpenAI API is not free. You will need to set up billing information there to be able to use OpenAI API. Alternatively, you can use models from HuggingFace Hub or other places. Check out my previous [blog post](#) and [video](#) on how to use other models.

```
import os
os.environ["OPENAI_API_KEY"] = "COPY AND PASTE YOUR API KEY HERE"
```



Sophia Yang

3.4K Followers

Ph.D. | Senior Data Scientist@Anaconda | Twitter: twitter.com/sophiamyang | YouTube: youtube.com/SophiaYangDS | Book Club: dsbookclub.github.io

Follow

Load documents

LangChain supports many many [Document Loaders](#) such as Notion, YouTube, and Figma. In this example, I'd like to chat with my PDF file. Thus, I used the PyPDFLoader to load my file. I'm actually using Chapter 1 of the [AI index report](#), which includes 55 pages, and I saved it in the materials directory of my Github [repo](#).

```
# Load document
from langchain.document_loaders import PyPDFLoader
loader = PyPDFLoader("materials/example.pdf")
documents = loader.load()
```

Method 1: load_qa_chain

`load_qa_chain` provides the most generic interface for answering questions. It loads a chain that you can do QA for your input documents and uses ALL of the text in the documents.

```
from langchain.llms import OpenAI
from langchain.chains.question_answering import load_qa_chain

chain = load_qa_chain(llm=OpenAI(), chain_type="map_reduce")
query = "How many AI publications?"
chain.run(input_documents=documents, question=query)
```

' The total number of AI publications has more than doubled since 2010, growing from 200,000 in 2010 to almost 500,000 in 2021.'

It also lets you do QA over a set of documents:

```
### For multiple documents
loaders = [...]
documents = []
for loader in loaders:
    documents.extend(loader.load())
```

😓 *But what if my document is super long that it exceeds the token limit?*

There are two ways to fix it:

Solution 1: Chain Type

The default `chain_type="stuff"` uses ALL of the text from the documents in the prompt. It actually doesn't work with our example because it exceeds the token limit and causes rate-limiting errors. That's why in this example, we had to use other chain types for example `"map_reduce"`. What are the other chain types?

- `map_reduce`: It separates texts into batches (as an example, you can define batch size in `llm=OpenAI(batch_size=5)`), feeds each batch with the question to LLM separately, and comes up with the final answer based on the answers from each batch.
- `refine`: It separates texts into batches, feeds the first batch to LLM, and feeds the answer and the second batch to LLM. It refines the answer by going through all the batches.
- `map-rerank`: It separates texts into batches, feeds each batch to LLM, returns a score of how fully it answers the question, and comes up with the final answer based on the high-scored answers from each batch.

Solution 2: RetrievalQA

One issue with using ALL of the text is that it can be very costly because you are feeding all the texts to OpenAI API and the API is charged by the number of tokens. A better solution is to retrieve relevant text chunks first and only use the relevant text chunks in the language model. I'm going to go through the details of RetrievalQA next.

Method 2: RetrievalQA

RetrievalQA chain actually uses `load_qa_chain` under the hood. We retrieve the most relevant chunk of text and feed those to the language model.

Here is how it works:

```
from langchain.chains import RetrievalQA
from langchain.indexes import VectorstoreIndexCreator
from langchain.text_splitter import CharacterTextSplitter
from langchain.embeddings import OpenAIEmbeddings
from langchain.vectorstores import Chroma

# split the documents into chunks
text_splitter = CharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
texts = text_splitter.split_documents(documents)
# select which embeddings we want to use
embeddings = OpenAIEmbeddings()
# create the vectorestore to use as the index
db = Chroma.from_documents(texts, embeddings)
# expose this index in a retriever interface
retriever = db.as_retriever(search_type="similarity", search_kwargs={"k":2})
# create a chain to answer questions
qa = RetrievalQA.from_chain_type(
    llm=OpenAI(), chain_type="stuff", retriever=retriever, return_source_documents=True
)
query = "How many AI publications in 2021?"
result = qa({"query": query})
```

In the result, we can see the answer and two source documents because we defined k as 2 meaning that we are only interested in getting two relevant text chunks.

Options:

There are various options for you to choose from in this process:

- embeddings: In the example, we used OpenAI Embeddings. But there are many other embedding options such as Cohere Embeddings, and HuggingFaceEmbeddings from specific models.
- TextSplitter: We used Character Text Splitter in the example where the text is split by a single character. You can also use different text splitters and different tokens mentioned in this [doc](#).
- VectorStore: We used Chroma as our vector database where we store our embedded text vectors. Other popular options are FAISS, Mulvus,

and Pinecone.

- **Retrievers:** We used a VectorStoreRetriever, which is backed by a VectorStore. To retrieve text, there are two search types you can choose: `search_type`: “similarity” or “mmr”. `search_type="similarity"` uses similarity search in the retriever object where it selects text chunk vectors that are most similar to the question vector.
`search_type="mmr"` uses the maximum marginal relevance search where it optimizes for similarity to query AND diversity among selected documents.
- **Chain Type:** same as method 1. You can also define the chain type as one of the four options: “stuff”, “map reduce”, “refine”, “map_rerank”.

Method 3: VectorstoreIndexCreator

VectorstoreIndexCreator is a wrapper around the above functionality. It is exactly the same under the hood, but just exposes a higher-level interface to let you get started in three lines of code:

Of course, you can also specify different options in this wrapper:

Method 4: ConversationalRetrievalChain

ConversationalRetrievalChain is very similar to method 2 RetrievalQA. It added an additional parameter `chat_history` to pass in chat history which can be used for follow-up questions.

ConversationalRetrievalChain = conversation memory + RetrievalQAChain

If you would like your language model to have a memory of the previous conversation, use this method. In my example below, I asked about the number of AI publications and got the result of 500,000. Then I asked the LLM to divide this number by 2. Since it has all the chat history, the model knows the number I was referring to is 500,000 and the result returned is 250,000.

Conclusion

Now you know four ways to do question answering with LLMs in LangChain. In summary, `load_qa_chain` uses all texts and accepts multiple documents; `RetrievalQA` uses `load_qa_chain` under the hood but retrieves relevant text chunks first; `VectorstoreIndexCreator` is the same as `RetrievalQA` with a higher-level interface; `ConversationalRetrievalChain` is useful when you want to pass in your chat history to the model.

Acknowledgment:

Thank you Harrison Chase for the guidance!

By [Sophia Yang](#) on April 8, 2023

Sophia Yang is a Senior Data Scientist. Connect with me on [LinkedIn](#), [Twitter](#), and [YouTube](#) and join the DS/ML [Book Club](#) ❤️

Langchain

OpenAI

Chatbots



More from Towards Data Science


Follow

Your home for data science. A Medium publication sharing concepts, ideas and codes.

Read more from Towards Data Science



[About](#) [Help](#) [Terms](#) [Privacy](#)





ServiceNow Admin Full Course | Learn ServiceNow Administration in 7.5 Hours| System Administration



SAASWITHSERVICENOW
42.5K subscribers

Subscribe

2.7K



Share

Save



200K views · 11 months ago · #SAASWITHSERVICENOW #ServiceNowSystemAdministrationTraining #ServiceNowJobs
If you want to support me then by me a coffee- <https://www.buymeacoffee.com/saasnow>

Please Note: This training has been prepared in Orlando version of ServiceNow. [Show more](#)

Skip navigation

You're signed out of YouTube

Sign in to like videos, comment, and subscribe.

Got it

SERVICENOW DEVELOPMENT FULL COURSE
9 HOURS

SERVICENOW INCIDENT MANAGEMENT
FOR ADMINS AND IT USERS
A COMPLETE TUTORIAL
1:12:47

#1 #ServiceNow #Incident Management | A Complete...
SAASWITHSERVICENOW
148K views · 4 years ago

SERVICENOW ADMIN FULL COURSE
7.5 HOURS (••)

Mix - SAASWITHSERVICENOW
More from this channel for you

servicenow FULL COURSE
IN 11 HOURS

ServiceNow Full Course in 2022 | 100 % Useful course | Contac...
Gautham Digital Learning
40K views · 8 months ago

SERVICENOW SYSTEM ADMINISTRATION TRAINING
#1 Introduction and Overview

Complete #ServiceNow System Administrator Training
SAASWITHSERVICENOW

SERVICENOW CMDB A COMPLETE TRAINING
Overview of CMDB

#1 What is CMDB | Overview of CMDB | ServiceNow CMDB...
SAASWITHSERVICENOW
76K views · 1 year ago

SERVICENOW FLOW DESIGNER TRAINING
Overview of Flow Designer

ServiceNow Flow Designer Training
SAASWITHSERVICENOW

UI Policy & Data Policy
What is the Difference?

UI Policy and Data Policy in ServiceNow | What is UI Policy...
SAASWITHSERVICENOW
15K views · 2 years ago

ServiceNow Users, Groups & Roles | ServiceNow Training |...
SAASWITHSERVICENOW
12K views · 11 months ago

#1 ServiceNow Developer Training | Preparation for...
SAASWITHSERVICENOW
117K views · 3 years ago

#3 #ServiceNow System Administration Training | List...
SAASWITHSERVICENOW
65K views · 2 years ago

#7 #ServiceNow System Administration Training |...
SAASWITHSERVICENOW
27K views · 2 years ago

Full CMDB In 3 Hours with Real World demonstration | ITOM |...
ServiceNow Helpdesk
55K views · 2 years ago

#5 #ServiceNow System Administration Training | Task...
SAASWITHSERVICENOW
49K views · 2 years ago

ServiceNow interview Questions and Answers for...
Gautham Digital Learning
14K views · 11 months ago

#1 #ServiceNow #Change #Management | A Complete...
SAASWITHSERVICENOW
111K views · 3 years ago

#6 #ServiceNow System Administration Training |...
SAASWITHSERVICENOW
40K views · 2 years ago



Enjoy unlimited DVR space

Record live TV & watch on your own time with YouTube TV. New users only. Terms apply. Cancel anytime.

No thanks

[Try it free](#)

You're signed out of YouTube

Sign in to like videos, comment, and subscribe.

Got it

48K views • 2 years ago

ServiceNow most asked 35 Interview questions with...

SKFacts and ITCareers

34K views • 11 months ago



Enjoy unlimited DVR space

Record live TV & watch on your own time with YouTube TV. New users only. Terms apply. Cancel anytime.

No thanks

[Try it free](#)

Task Requirements:

Step Number 2:

```
from langchain.document_loaders import UnstructuredPDFLoader, OnlinePDFLoader

from langchain.text_splitter import RecursiveCharacterTextSplitter
```

Load your data

In [2]:

```
loader = UnstructuredPDFLoader("../data/field-guide-to-data-science.pdf")

# loader =
OnlinePDFLoader("https://wolfpaulus.com/wp-content/uploads/2017/05/field-guide-to-data-science.pdf")
```

In [3]:

```
data = loader.load()
```

In [4]:

```
print(f'You have {len(data)} document(s) in your data')

print(f'There are {len(data[0].page_content)} characters in your document')
```

You have 1 document(s) in your data

There are 176584 characters in your document

Chunk your data up into smaller documents

In [5]:

Task Requirements:

```
text_splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
```

```
texts = text_splitter.split_documents(data)
```

In [6]:

```
print (f'Now you have {len(texts)} documents')
```

Now you have 228 documents

Create embeddings of your documents to get ready for semantic search

In [9]:

```
from langchain.vectorstores import Chroma, Pinecone
```

```
from langchain.embeddings.openai import OpenAIEmbeddings
```

```
import pinecone
```

In [24]:

```
OPENAI_API_KEY = '...'
```

```
PINECONE_API_KEY = '...'
```

```
PINECONE_API_ENV = 'us-east1-gcp'
```

In [11]:

```
embeddings = OpenAIEmbeddings(openai_api_key=OPENAI_API_KEY)
```

In [12]:

```
# initialize pinecone
```

Task Requirements:

```
pinecone.init(  
  
    api_key=PINECONE_API_KEY, # find at app.pinecone.io  
  
    environment=PINECONE_API_ENV # next to api key in console  
  
)  
  
index_name = "langchain2"
```

In [13]:

```
docsearch = Pinecone.from_texts([t.page_content for t in texts], embeddings,  
index_name=index_name)
```

In [14]:

```
query = "What are examples of good data science teams?"  
  
docs = docsearch.similarity_search(query, include_metadata=True)
```

Query those docs to get your answer back

In [16]:

```
from langchain.llms import OpenAI  
  
from langchain.chains.question_answering import load_qa_chain
```

In [17]:

```
llm = OpenAI(temperature=0, openai_api_key=OPENAI_API_KEY)  
  
chain = load_qa_chain(llm, chain_type="stuff")
```

In [22]:

Task Requirements:

```
query = "What is the collect stage of data maturity?"
```

```
docs = docsearch.similarity_search(query, include_metadata=True)
```

In [23]:

```
chain.run(input_documents=docs, question=query)
```

Out[23]:

```
' The collect stage of data maturity focuses on collecting internal or external datasets. Examples include gathering sales records and corresponding weather data.'
```

In []:

LangChain QA

All code comes from [LangChain docs](#).

In []:

```
!pip install langchain openai chromadb tiktoken pypdf
```

In []:

```
import os
os.environ["OPENAI_API_KEY"] = ""
```

In []:

```
from langchain.chains import RetrievalQA
from langchain.llms import OpenAI
from langchain.document_loaders import TextLoader
from langchain.document_loaders import PyPDFLoader
from langchain.indexes import VectorstoreIndexCreator
from langchain.text_splitter import CharacterTextSplitter
from langchain.embeddings import OpenAIEmbeddings
from langchain.vectorstores import Chroma
```

In []:

```
llm = OpenAI()
print(llm("tell me a joke"))
```

Task Requirements:

load_qa_chain

Loads a chain that you can use to do QA over a set of documents, but it uses ALL of those documents.

chain_type="stuff" will not work because the number of tokens exceeds the limit. We can try other chain types like "map_reduce".

In []:

```
from langchain.chains.question_answering import load_qa_chain
```

```
# load document
```

```
loader = PyPDFLoader("materials/example.pdf")
```

```
documents = loader.load()
```

```
### For multiple documents
```

```
# loaders = [...]
```

```
# documents = []
```

```
# for loader in loaders:
```

```
#     documents.extend(loader.load())
```

```
chain = load_qa_chain(llm=OpenAI(), chain_type="map_reduce")
```

```
query = "what is the total number of AI publications?"
```

```
chain.run(input_documents=documents, question=query)
```

RetrievalQA

RetrievalQA chain uses load_qa_chain under the hood. We retrieve the most relevant chunk of text and feed those to the language model.

Options:

- [embeddings](#)
- [TextSplitter](#)
- [VectorStore](#)
- [Retrievers](#)
 - [search_type](#): "similarity" or "mmr"
- [Chain Type](#): "stuff", "map reduce", "refine", "map_rerank"

In []:

```
# load document
```

```
loader = PyPDFLoader("materials/example.pdf")
```

```
documents = loader.load()
```

```
# split the documents into chunks
```

```
text_splitter = CharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
```

```
texts = text_splitter.split_documents(documents)
```

```
# select which embeddings we want to use
```

```
embeddings = OpenAIEmbeddings()
```

```
# create the vectorestore to use as the index
```

Task Requirements:

```
db = Chroma.from_documents(texts, embeddings)
# expose this index in a retriever interface
retriever = db.as_retriever(search_type="similarity", search_kwargs={"k":2})
# create a chain to answer questions
qa = RetrievalQA.from_chain_type(
    llm=OpenAI(), chain_type="stuff", retriever=retriever, return_source_documents=True)
query = "what is the total number of AI publications?"
result = qa({"query": query})
```

In []:

```
retriever.get_relevant_documents(query)
```

In []:

```
result
```

VectorstoreIndexCreator

VectorstoreIndexCreator is a wrapper for the above logic.

Source:

- https://python.langchain.com/en/latest/modules/chains/getting_started.html
- <https://github.com/hwchase17/langchain/blob/master/langchain/indexes/vectorstore.py#L21-L74>

In []:

```
index = VectorstoreIndexCreator(
    # split the documents into chunks
    text_splitter=CharacterTextSplitter(chunk_size=1000, chunk_overlap=0),
    # select which embeddings we want to use
    embedding=OpenAIEmbeddings(),
    # use Chroma as the vectorstore to index and search embeddings
    vectorstore_cls=Chroma
).from_loaders([loader])
query = "what is the total number of AI publications?"
index.query(llm=OpenAI(), question=query, chain_type="stuff")
```

ConversationalRetrievalChain

conversation memory + RetrievalQAChain

Allow for passing in chat history which can be used for follow up questions.

Source:

https://python.langchain.com/en/latest/modules/chains/index_examples/chat_vector_db.html

Task Requirements:

```
from langchain.chains import ConversationalRetrievalChain
```

In []:

```
# load document
loader = PyPDFLoader("materials/example.pdf")
documents = loader.load()
# split the documents into chunks
text_splitter = CharacterTextSplitter(chunk_size=1000, chunk_overlap=0)
texts = text_splitter.split_documents(documents)
# select which embeddings we want to use
embeddings = OpenAIEmbeddings()
# create the vectorestore to use as the index
db = Chroma.from_documents(texts, embeddings)
# expose this index in a retriever interface
retriever = db.as_retriever(search_type="similarity", search_kwargs={"k":2})
# create a chain to answer questions
qa = ConversationalRetrievalChain.from_llm(OpenAI(), retriever)
chat_history = []
query = "what is the total number of AI publications?"
result = qa({"question": query, "chat_history": chat_history})
```

In []:

```
chat_history = []
query = "what is the total number of AI publications?"
result = qa({"question": query, "chat_history": chat_history})
```

In []:

```
result["answer"]
```

In []:

```
chat_history = [(query, result["answer"])]
query = "What is this number divided by 2?"
result = qa({"question": query, "chat_history": chat_history})
```

In []:

```
chat_history
```

In []:

```
result['answer']
```

In []:

Step Number 3:

A- PineCone package

Task Requirements:

This page provides installation instructions, usage examples, and a reference for the Pinecone Node.JS client.



Warning

This is a public preview ("Beta") client. Test thoroughly before using this client for production workloads. No SLAs or technical support commitments are provided for this client. Expect potential breaking changes in future releases.

Getting Started

Installation

Use the following shell command to install the Node.JS client for use with Node.JS versions 17 and above:

Shell

```
npm install @pinecone-database/pinecone
```

Alternatively, you can install Pinecone with Yarn:

Shell

```
yarn add @pinecone-database/pinecone
```

Usage

Initialize the client

To initialize the client, instantiate the `PineconeClient` class and call the `init` method. The `init` method takes an object with the `apiKey` and `environment` properties:

JavaScript

```
import { PineconeClient } from  
"@pinecone-database/pinecone";
```

```
const pinecone = new PineconeClient();  
await pinecone.init({
```

Task Requirements:

```
environment: "YOUR_ENVIRONMENT",  
apiKey: "YOUR_API_KEY",  
});
```

Create index

The following example creates an index without a metadata configuration. By default, Pinecone indexes all metadata.

JavaScript

```
await pinecone.createIndex({  
  createRequest: {  
    name: "example-index",  
    dimension: 1024,  
  },  
});
```

The following example creates an index that only indexes the "color" metadata field. Queries against this index cannot filter based on any other metadata field.

JavaScript

```
await pinecone.createIndex({  
  createRequest: {  
    name: "example-index-2",  
    dimension: 1024,  
    metadata_config: {  
      indexed: ["color"],  
    },  
  },  
});
```

List indexes

The following example logs all indexes in your project.

JavaScript

```
const indexesList = await pinecone.listIndexes();
```

Describe index

The following example logs information about the index example-index.

Task Requirements:

JavaScript

```
const indexDescription = await pinecone.describeIndex({
  indexName: "example-index",
});
```

Delete index

The following example deletes example-index.

JavaScript

```
await pinecone.deleteIndex({
  indexName: "example-index",
});
```

Scale replicas

The following example sets the number of replicas and pod type for example-index.

JavaScript

```
await pinecone.configureIndex({
  indexName: "example-index",
  patchRequest: {
    replicas: 2,
    podType: "p2",
  },
});
```

Describe index statistics

The following example returns statistics about the index example-index.

JavaScript

```
const index = pinecone.Index("example-index");
const indexStats = index.describeIndexStats({
  describeIndexStatsRequest: {
    filter: {},
  },
});
```

Upsert vectors

The following example upserts vectors to example-index.

Task Requirements:

JavaScript

```
const index = pinecone.Index("example-index");
const upsertRequest = {
  vectors: [
    {
      id: "vec1",
      values: [0.1, 0.2, 0.3, 0.4],
      metadata: {
        genre: "drama",
      },
    },
    {
      id: "vec2",
      values: [0.2, 0.3, 0.4, 0.5],
      metadata: {
        genre: "action",
      },
    },
  ],
  namespace: "example-namespace",
};
const upsertResponse = await index.upsert({ upsertRequest
});
```

Query an index

The following example queries the index example-index with metadata filtering.

JavaScript

```
const index = pinecone.Index("example-index");
const queryRequest = {
  vector: [0.1, 0.2, 0.3, 0.4],
  topK: 10,
  includeValues: true,
  includeMetadata: true,
  filter: {
    genre: { $in: ["comedy", "documentary", "drama"] },
  },
};
```

Task Requirements:

```
namespace: "example-namespace",  
};  
const queryResponse = await index.query({ queryRequest });
```

Delete vectors

The following example deletes vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");  
await index.delete1({  
  ids: ["vec1", "vec2"],  
  namespace: "example-namespace",  
});
```

Fetch vectors

The following example fetches vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");  
const fetchResponse = await index.fetch({  
  ids: ["vec1", "vec2"],  
  namespace: "example-namespace",  
});
```

Update vectors

The following example updates vectors by ID.

JavaScript

```
const index = pinecone.Index("example-index");  
const updateRequest = {  
  id: "vec1",  
  values: [0.1, 0.2, 0.3, 0.4],  
  setMetadata: { genre: "drama" },  
  namespace: "example-namespace",  
};
```

```
const updateResponse = await index.update({  
  updateRequest });
```

Create collection

The following example creates the collection example-collection from example-index.

Task Requirements:

JavaScript

```
const createCollectionRequest = {  
  name: "example-collection",  
  source: "example-index",  
};
```

```
await pinecone.createCollection({  
  createCollectionRequest,  
});
```

List collections

The following example returns a list of the collections in the current project.

JavaScript

```
const collectionsList = await pinecone.listCollections();
```

Describe a collection

The following example returns a description of the collection example-collection.

JavaScript

```
const collectionDescription = await  
pinecone.describeCollection({  
  collectionName: "example-collection",  
});
```

Delete a collection

The following example deletes the collection example-collection.

JavaScript

```
await pinecone.deleteCollection({  
  collectionName: "example-collection",  
});
```

Reference

For the REST API or other clients, see the [API reference](#).

Task Requirements:

init()

pinecone.init(configuration: PineconeClientConfiguration)

Initialize the Pinecone client.

Parameters **Type** **Description**

configuration **PineconeClientConfiguration** The configuration for the Pinecone client.

Types

PineconeClientConfiguration

Parameters **Type** **Description**

apiKey **string** The API key for the Pinecone service.

environment **string** The cloud environment of your Pinecone project.

Example:

JavaScript

```
import { PineconeClient } from
"@pinecone-database/pinecone";
const pinecone = new PineconeClient();
await pinecone.init({
  apiKey: "YOUR_API_KEY",
  environment: "YOUR_ENVIRONMENT",
});
configureIndex()
pinecone.configure_index(indexName: string,
patchRequest?: PatchRequest)
```

Configure an index to change pod type and number of replicas.

Parameters **Type** **Description**

requestParameters **ConfigureIndexRequest** **Index**
configuration parameters.

Types

ConfigureIndexRequest

Parameters **Type** **Description**

indexName **string** The name of the index.

Task Requirements:

patchRequest **PatchRequest** (Optional) Patch request parameters.

PatchRequest

Parameters	Type	Description
------------	------	-------------

replicas	number	(Optional) The number of replicas to configure for this index.
-----------------	---------------	--

podType	string	(Optional) The new pod type for the index. One of s1, p1, or p2 appended with . and one of x1, x2, x4, or x8.
----------------	---------------	---

Example:

JavaScript

```
const newNumberOfReplicas = 4;
const newPodType = "s1.x4";
await pinecone.configureIndex({
  indexName: "example-index",
  patchRequest: {
    replicas: newNumberOfReplicas,
    podType: newPodType,
  },
});
createCollection()
pinecone.createCollection(requestParameters:
CreateCollectionOperationRequest)
```

Create a collection from an index.

Parameters	Type	Description
------------	------	-------------

requestParameters	CreateCollectionOperationRequest
--------------------------	---

Create collection operation wrapper

Types

CreateCollectionOperationRequest

Parameters	Type	Description
------------	------	-------------

createCollectionRequest	CreateCollectionRequest
--------------------------------	--------------------------------

Collection request parameters.

CreateCollectionRequest

Parameters	Type	Description
------------	------	-------------

name	string	The name of the collection to be created.
-------------	---------------	---

Task Requirements:

source **string** The name of the source index to be used as the source for the collection.

Example:

JavaScript

```
await pinecone.createCollection({
  createCollectionRequest: {
    name: "example-collection",
    source: "example-index",
  },
});
createIndex()
pinecone.createIndex(requestParameters?:
CreateIndexRequest)
```

Create an index.

Parameters **Type** **Description**

requestParameters **CreateIndexRequest** **Create index operation wrapper**

Types

CreateIndexRequest

Parameters **Type** **Description**

createRequest **CreateRequest** **Create index request parameters**

CreateRequest

Parameters **Type** **Description**

name **str** The name of the index to be created. The maximum length is 45 characters.

dimension **integer** The dimensions of the vectors to be inserted in the index.

metric **str** (Optional) The distance metric to be used for similarity search: 'euclidean', 'cosine', or 'dotproduct'.

pods **int** (Optional) The number of pods for the index to use, including replicas.

replicas **int** (Optional) The number of replicas.

pod_type **str** (Optional) The new pod type for the index. One of s1, p1, or p2 appended with . and one of x1, x2, x4, or x8.

Task Requirements:

metadata_config object (Optional) Configuration for the behavior of Pinecone's internal metadata index. By default, all metadata is indexed; when metadata_config is present, only specified metadata fields are indexed. To specify metadata fields to index, provide a JSON object of the following form: {"indexed": ["example_metadata_field"]}

source_collection str (Optional) The name of the collection to create an index from.

Example:

JavaScript

// The following example creates an index without a metadata configuration. By default, Pinecone indexes all metadata.

```
await pinecone.createIndex({
  createRequest: {
    name: "pinecone-index",
    dimension: 1024,
  },
});
```

// The following example creates an index that only indexes the 'color' metadata field. Queries against this index cannot filter based on any other metadata field.

```
await pinecone.createIndex({
  createRequest: {
    name: "example-index-2",
    dimension: 1024,
    metadata_config: {
      indexed: ["color"],
    },
  },
});
deleteCollection()
pinecone.deleteCollection(requestParameters:
DeleteCollectionRequest)
```

Delete an existing collection.

Task Requirements:

Types

Parameters

Type	Description
------	-------------

requestParameters	DeleteCollectionRequest	Delete
--------------------------	--------------------------------	---------------

collection request parameters

DeleteCollectionRequest

Parameters

Type	Description
------	-------------

collectionName	string	The name of the collection to delete.
-----------------------	---------------	---------------------------------------

Example:

JavaScript

```
await pinecone.deleteCollection({
  collectionName: "example-collection",
});
deleteIndex()
pinecone.deleteIndex(requestParameters:
DeleteIndexRequest)
```

Delete an index.

Types

Parameters

Type	Description
------	-------------

requestParameters	DeleteIndexRequest	Delete index
--------------------------	---------------------------	---------------------

request parameters

DeleteIndexRequest

Parameters

Type	Description
------	-------------

indexName	string	The name of the index to delete.
------------------	---------------	----------------------------------

Example:

JavaScript

```
await pinecone.deleteIndex({
  indexName: "example-index",
});
describeCollection()
pinecone.describeCollection(requestParameters:
DescribeCollectionRequest)
```

Get a description of a collection.

Task Requirements:

Types

Parameters Type Description

requestParameters **DescribeCollectionRequest** Describe collection request parameters

DescribeCollectionRequest

Parameters Type Description

collectionName **string** The name of the collection.

Example:

JavaScript

```
const collectionDescription = await
pinecone.describeCollection({
  collectionName: "example-collection",
});
```

Return:

collectionMeta : object Configuration information and deployment status of the collection.

name : string The name of the collection.

size: integer The size of the collection in bytes.

status: string The status of the collection.

describeIndex()

pinecone.describeIndex(requestParameters: DescribeIndexRequest)

Get a description of an index.

Types

Parameters Type Description

requestParameters **DescribeIndexRequest** Describe index request parameters

DescribeIndexRequest

Parameters Type Description

indexName string The name of the index.

Types

Returns:

database : object

Task Requirements:

name : string The name of the index.

dimension : integer The dimensions of the vectors to be inserted in the index.

metric : string The distance metric used for similarity search: 'euclidean', 'cosine', or 'dotproduct'.

pods : integer The number of pods the index uses, including replicas.

replicas : integer The number of replicas.

pod_type : string The pod type for the index. One of s1, p1, or p2 appended with . and one of x1, x2, x4, or x8.

metadata_config: object Configuration for the behavior of Pinecone's internal metadata index. By default, all metadata is indexed; when metadata_config is present, only specified metadata fields are indexed. To specify metadata fields to index, provide a JSON object of the following form:

```
{"indexed": ["example_metadata_field"]}
```

status : object

ready : boolean Whether the index is ready to serve queries.

state : string One of Initializing, ScalingUp, ScalingDown, Terminating, or Ready.

Example:

JavaScript

```
const indexDescription = await pinecone.describeIndex({
  indexName: "example-index",
});
listCollections
pinecone.listCollections()
```

Return a list of the collections in your project.

Example:

JavaScript

```
const collections = await pinecone.listCollections();
Returns:
```

array of strings The names of the collections in your project.

Task Requirements:

listIndexes

pinecone.listIndexes()

Return a list of your Pinecone indexes.

Returns:

array of strings The names of the indexes in your project.

Example:

JavaScript

```
const indexesList = await pinecone.listIndexes();  
Index()  
pinecone.Index(indexName: string)
```

Construct an Index object.

Parameters

Type	Description
------	-------------

indexName	string	The name of the index.
-----------	--------	------------------------

Example:

JavaScript

```
const index = pinecone.Index("example-index");  
Index.delete1()  
index.delete(requestParameters: Delete1Request)
```

Delete items by their ID from a single namespace.

Parameters

Type	Description
------	-------------

requestParameters	Delete1Request	Delete request parameters
-------------------	----------------	---------------------------

Types

Delete1Request

Parameters

Type	Description
------	-------------

ids	Array (Optional)	The IDs of the items to delete.
-----	------------------	---------------------------------

deleteAll	boolean	(Optional) Indicates that all vectors in the index namespace should be deleted.
-----------	---------	---

Task Requirements:

namespacestr (Optional) The namespace to delete vectors from, if applicable.

Types

Example:

JavaScript

```
await index.delete1({
  ids: ["example-id-1", "example-id-2"],
  namespace: "example-namespace",
});
Index.describeIndexStats()
index.describeIndexStats(requestParameters:
DescribeIndexStatsOperationRequest)
```

Returns statistics about the index's contents, including the vector count per namespace and the number of dimensions.

Parameters	Type	Description
------------	------	-------------

requestParameters	DescribeIndexStatsOperationRequest	Describe index stats request wrapper
-------------------	------------------------------------	--------------------------------------

Types

DescribeIndexStatsOperationRequest

Parameters	Type	Description
------------	------	-------------

describeIndexStatsRequest	DescribeIndexStatsRequest	Describe index stats request parameters
---------------------------	---------------------------	---

DescribeIndexStatsRequest

parameter	Type	Description
-----------	------	-------------

filter	object	(Optional) A metadata filter expression.
--------	--------	--

Returns:

namespaces : object A mapping for each namespace in the index from the namespace name to a summary of its contents. If a metadata filter expression is present, the summary will reflect only vectors matching that expression.

dimension : int64 The dimension of the indexed vectors.

indexFullness : float The fullness of the index, regardless of whether a metadata filter expression was passed. The granularity of this metric is 10%.

Task Requirements:

totalVectorCount : int64 The total number of vectors in the index.

Example:

JavaScript

```
const indexStats = await index.describeIndexStats({
  describeIndexStatsRequest: {},
});
```

Read more about filtering for more detail.

Index.fetch()

index.fetch(requestParameters: FetchRequest)

The Fetch operation looks up and returns vectors, by ID, from a single namespace. The returned vectors include the vector data and metadata.

Parameters	Type	Description
------------	------	-------------

requestParameters	FetchRequest	Fetch request parameters
-------------------	--------------	--------------------------

Types

FetchRequest

Parameters	Type	Description
------------	------	-------------

ids	Array	The vector IDs to fetch. Does not accept values containing spaces.
-----	-------	--

namespace	string	(Optional) The namespace containing the vectors.
-----------	--------	--

Returns:

vectors : object Contains the vectors.

namespace : string The namespace of the vectors.

Example:

JavaScript

```
const fetchResponse = await index.fetch({
  ids: ["example-id-1", "example-id-2"],
  namespace: "example-namespace",
});
```

Task Requirements:

Index.query()

index.query(requestParameters: QueryOperationRequest)

Search a namespace using a query vector. Retrieves the ids of the most similar items in a namespace, along with their similarity scores.

Parameters

Type	Description
------	-------------

requestParameters	QueryOperationRequest	The query operation request wrapper.
--------------------------	------------------------------	--------------------------------------

Types

Parameters

Type	Description
------	-------------

queryRequest	QueryRequest	The query operation request.
---------------------	---------------------	------------------------------

QueryRequest

Parameter

Type	Description
------	-------------

namespacestring	(Optional)	The namespace to query.
------------------------	------------	-------------------------

topK	number	The number of results to return for each query.
-------------	--------	---

filter	object	(Optional) The filter to apply. You can use vector metadata to limit your search. See https://www.pinecone.io/docs/metadata-filtering/ .
---------------	--------	---

includeValues	boolean	(Optional) Indicates whether vector values are included in the response. Defaults to false.
----------------------	---------	---

includeMetadata	boolean	(Optional) Indicates whether metadata is included in the response as well as the ids. Defaults to false.
------------------------	---------	--

vector	Array	(Optional) The query vector. This should be the same length as the dimension of the index being queried. Each query() request can contain only one of the parameters id or vector.
---------------	-------	--

id	string	(Optional) The unique ID of the vector to be used as a query vector. Each query() request can contain only one of the parameters vector or id.
-----------	--------	--

Example:

JavaScript

```
const queryResponse = await index.query({
  queryRequest: {
```

Task Requirements:

```
namespace: "example-namespace",
topK: 10,
filter: {
  genre: { $in: ["comedy", "documentary", "drama"] },
},
includeValues: true,
includeMetadata: true,
vector: [0.1, 0.2, 0.3, 0.4],
},
});
Index.update()
index.update(requestParameters: UpdateOperationRequest)
```

Updates vectors in a namespace. If a value is included, it will overwrite the previous value.

If setMetadata is included in the updateRequest, the values of the fields specified in it will be added or overwrite the previous value.

Parameters Type Description

requestParameters	UpdateOperationRequest	The update operation wrapper
-------------------	------------------------	------------------------------

Types

UpdateOperationRequest

Parameters Type Description

updateRequest	UpdateRequest	The update request.
---------------	---------------	---------------------

UpdateRequest

Parameter Type Description

id	string	The vector's unique ID.
----	--------	-------------------------

values	Array (Optional)	Vector data.
--------	------------------	--------------

setMetadata	object	(Optional) Metadata to set for the vector.
-------------	--------	--

namespace	string	(Optional) The namespace containing the vector.
-----------	--------	---

Example:

JavaScript

```
const updateResponse = await index.update({
  updatedRequest: {
```


Task Requirements:

```
id: "vec1",
values: [0.1, 0.2, 0.3, 0.4],
setMetadata: {
  genre: "drama",
},
namespace: "example-namespace",
},
});
Index.upsert()
index.upsert(requestParameters: UpsertOperationRequest)
```

Writes vectors into a namespace. If a new value is upserted for an existing vector ID, it will overwrite the previous value.

Parameters Type Description

requestParameters	UpsertOperationRequest	Upsert operation wrapper
-------------------	------------------------	--------------------------

Types

UpsertOperationRequest

Parameters Type Description

upsertRequest	UpsertRequest	The upsert request.
---------------	---------------	---------------------

UpsertRequest

Parameter	Type	Description
-----------	------	-------------

vectors	Array	An array containing the vectors to upsert.
---------	-------	--

Recommended batch limit is 100 vectors.

id (str) - The vector's unique id.

values ([float]) - The vector data.

metadata (object) - (Optional) Metadata for the vector. |

| namespace | string | (Optional) The namespace name to upsert vectors. |

Vector

Parameter Type Description

id	string	The vector's unique ID.
----	--------	-------------------------

values	Array	Vector data.
--------	-------	--------------

metadata	object	(Optional) Metadata for the vector.
----------	--------	-------------------------------------

Returns:

upsertedCount : int64 The number of vectors upserted.

Example:

Task Requirements:

JavaScript

```
const upsertResponse = await index.upsert({
  upsertRequest: {
    vectors: [
      {
        id: "vec1",
        values: [0.1, 0.2, 0.3, 0.4],
        metadata: {
          genre: "drama",
        },
      },
      {
        id: "vec2",
        values: [0.1, 0.2, 0.3, 0.4],
        metadata: {
          genre: "comedy",
        },
      },
    ],
    namespace: "example-namespace",
  },
});
```

b- Open AI

OpenAI Node.js Library

The OpenAI Node.js library provides convenient access to the OpenAI API from Node.js applications. Most of the code in this library is generated from our [OpenAPI specification](#).

Important note: this library is meant for server-side usage only, as using it in client-side browser code will expose your secret API key.

See [here](#) for more details.

Task Requirements:

Installation

```
$ npm install openai
```

Usage

The library needs to be configured with your account's secret key, which is available on the [website](#). We recommend setting it as an environment variable. Here's an example of initializing the library with the API key loaded from an environment variable and creating a completion:

```
const { Configuration, OpenAIApi } = require("openai");
```

```
const configuration = new Configuration({
```

```
  apiKey: process.env.OPENAI_API_KEY,
```

```
});
```

```
const openai = new OpenAIApi(configuration);
```

```
const completion = await openai.createCompletion({
```

```
  model: "text-davinci-003",
```

```
  prompt: "Hello world",
```

```
});
```

Task Requirements:

```
console.log(completion.data.choices[0].text);
```

Check out the [full API documentation](#) for examples of all the available functions.

Request options

All of the available API request functions additionally contain an optional final parameter where you can pass custom [axios request options](#), for example:

```
const completion = await openai.createCompletion(  
  
  {  
  
    model: "text-davinci-003",  
  
    prompt: "Hello world",  
  
  },  
  
  {  
  
    timeout: 1000,  
  
    headers: {  
  
      "Example-Header": "example",  
  
    },  
  
  }  
);
```

Task Requirements:

Error handling

API requests can potentially return errors due to invalid inputs or other issues. These errors can be handled with a `try...catch` statement, and the error details can be found in either `error.response` or `error.message`:

```
try {  
  
    const completion = await openai.createCompletion({  
  
        model: "text-davinci-003",  
  
        prompt: "Hello world",  
  
    });  
  
    console.log(completion.data.choices[0].text);  
  
} catch (error) {  
  
    if (error.response) {  
  
        console.log(error.response.status);  
  
        console.log(error.response.data);  
  
    } else {  
  
        console.log(error.message);  
  
    }  
}
```

Streaming completions

Task Requirements:

Streaming completions (`stream=true`) are not natively supported in this package yet, but [a workaround exists](#) if needed.

Upgrade guide

All breaking changes for major version releases are listed below.

3.0.0

- The function signature of `createCompletion(engineId, params)` changed to `createCompletion(params)`. The value previously passed in as the `engineId` argument should now be passed in as `model` in the `params` object (e.g. `createCompletion({ model: "text-davinci-003", ... })`)
- Replace any `createCompletionFromModel(params)` calls with `createCompletion(params)`

Thanks

Thank you to [ceifa](#) for creating and maintaining the original unofficial [openai](#) npm package before we released this official library! [ceifa's](#) original package has been renamed to [gpt-x](#).

Keywords

openai

open

ai

gpt-3

Gpt3

Task Requirements:

C-Langchain

Welcome to LangChain

LangChain is a framework for developing applications powered by language models. We believe that the most powerful and differentiated applications will not only call out to a language model via an api, but will also:

- *Be data-aware*: connect a language model to other sources of data
- *Be agentic*: allow a language model to interact with its environment

The LangChain framework is designed with the above principles in mind.

Getting Started

Checkout the below guide for a walkthrough of how to get started using LangChain to create an Language Model application.

- [Getting Started Documentation](#)

Components

There are several main modules that LangChain provides support for. For each module we provide some examples to get started and get familiar with some of the concepts. These modules are, in increasing order of complexity:

- **Prompts**: This includes prompt management, prompt optimization, and prompt serialization.
- **LLMs**: This includes a generic interface for all LLMs, and common utilities for working with LLMs.
- **Indexes**: This includes patterns and functionality for structuring your own text data so it can interact with language models (including embeddings, vectorstores, text splitters, retrievers, etc).
- **Memory**: Memory is the concept of persisting state between calls of a chain/agent. LangChain provides a standard interface for memory, a collection of memory implementations, and examples of chains/agents that use memory.
- **Chains**: Chains go beyond just a single LLM call, and are sequences of calls (whether to an LLM or a different utility). LangChain provides a standard interface for chains, lots of integrations with other tools, and end-to-end chains for common applications.
- **Agents**: Agents involve an LLM making decisions about which Actions to take, taking that Action, seeing an Observation, and repeating that until

Task Requirements:

done. LangChain provides a standard interface for agents, a selection of agents to choose from, and examples of end to end agents.

Reference Docs

All of LangChain's reference documentation, in one place. Full documentation on all methods and classes.

Production

As you move from prototyping into production, we're developing resources to help you do so. These including:

- **Deployment:** resources on how to deploy your end application.
- **Tracing:** resources on how to use tracing to log and debug your applications.

Additional Resources

Additional collection of resources we think may be useful as you develop your application!

- [LangChainHub](#): The LangChainHub is a place to share and explore other prompts, chains, and agents.
- [Discord](#): Join us on our Discord to discuss all things LangChain!
- [Production Support](#): As you move your LangChains into production, we'd love to offer more comprehensive support. Please fill out this form and we'll set up a dedicated support Slack channel